


Legged Robots Practical: Project 2

16.11.2021

Plan


- **W9 16.11.2021:** Quadruped Central Pattern Generators
+ Deep Reinforcement Learning
 - **W10 23.11.2021:**
 - **W11 30.11.2021:**
 - **W12 07.12.2021:**
 - **W13 14.12.2021:** Exam
 - **W14 21.12.2021:** Competition (Details to come)
- [Report 2 - Quadruped] - 30% of course grade**
- 

Quadruped Locomotion with Central Pattern Generators and Deep Reinforcement Learning


Legged Robots

Part 1: Central Pattern Generators

Trot




Bound





Pace






Walk




Part 2: Deep Reinforcement Learning






Quadruped Model Reference Frame




Quadruped Model Leg References



Quadruped Model Joint References



Joint angles \leftrightarrow Cartesian space (in leg frame)



$$p = f(q)$$

Forward kinematics

$$q = f^{-1}(p)$$

Inverse kinematics


$$\dot{p} = v = J(q)\dot{q}$$

Foot linear velocity

$$\tau = J^T(q)F$$

Map desired end effector force to torques

Joint angles \leftrightarrow Cartesian space (leg frame control)



$$p = f(q)$$

Forward kinematics

$$q = f^{-1}(p)$$

Inverse kinematics

$$\dot{p} = v = J(q)\dot{q}$$

Foot linear velocity

$$\tau = J^T(q)F$$

Map desired end effector force to torques

$$\tau_{joint} = K_{p,joint}(q_d - q) + K_{d,joint}(\dot{q}_d - \dot{q})$$

Joint PD

$$\tau_{Cartesian} = J^T(q)[K_{p,Cartesian}(p_d - p) + K_{d,Cartesian}(v_d - v)]$$

Cartesian PD

$$\tau_{final} = \tau_{joint} + \tau_{Cartesian}$$

Contributions from both joint PD and Cartesian PD

Central Pattern Generators: Review

100%


From Lecture 6

Descending modulation

Spinal cord

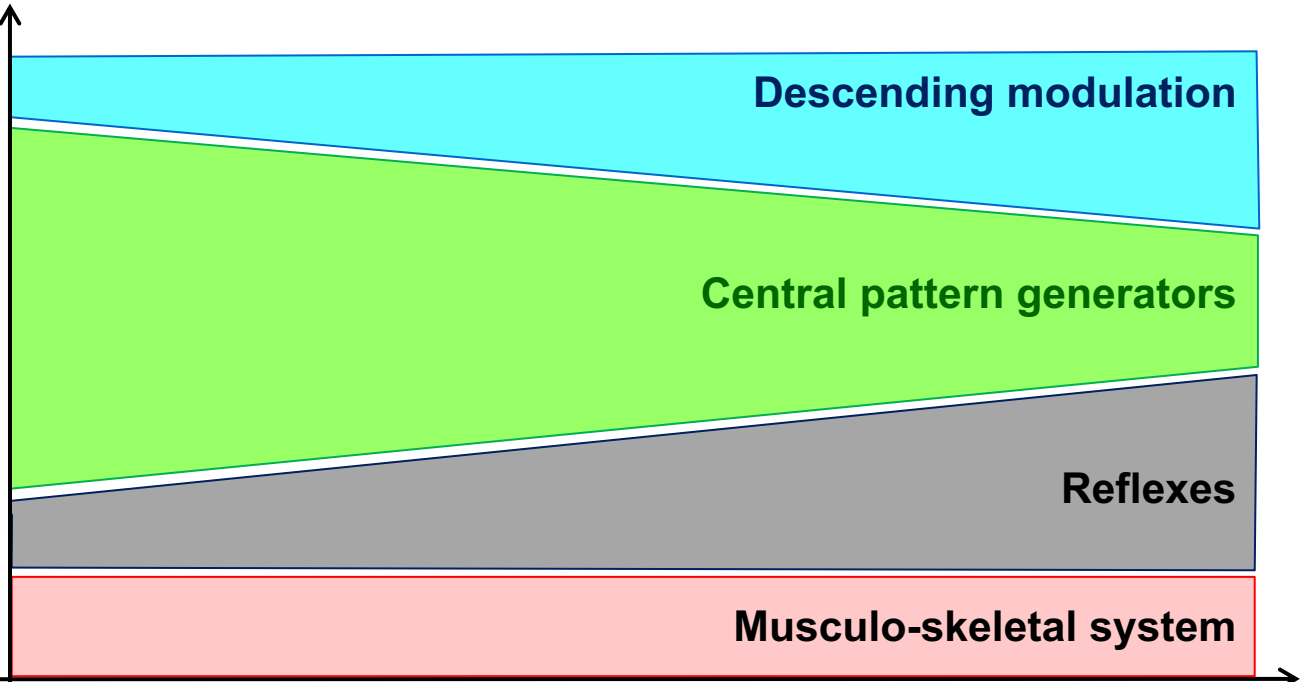
Reflexes

Central pattern generators



Musculoskeletal system

Respective Role
in motor control



“Complexity” of animal species

Ryczko, Simon, Ijspeert,
Trends in Neuroscience,
2020

Minassian et al
Neuroscientist
2017



lamprey



salamander



cat



human

100%


From Lecture 6

Descending modulation

Spinal cord

Reflexes

Central pattern generators



Musculoskeletal system

Respective Role
in motor control



lamprey




salamander



cat



human



Modeling the CPG with coupled oscillators

A segmental oscillator is modeled as an amplitude-controlled phase oscillator as used in (Cohen, Holmes and Rand 1982, Kopell, Ermentrout, and Williams 1990) :

Phase:

$$\dot{\theta}_i = 2\pi\nu_i + \sum_j r_j w_{ij} \sin(\theta_j - \theta_i - \phi_{ij})$$

Amplitude:

$$\dot{r}_i = a_i \left(\frac{a_i}{4} (R_i - r_i) - \dot{r}_i \right)$$



Output:

$$x_i = r_i (1 + \cos(\theta_i))$$


Setpoints:

$$\varphi_i = x_i - x_{N+i} \quad \text{for the axial motors}$$

$$\varphi_i = f(\theta_i) \quad \text{for the (rotational) limb motors}$$



From Lecture 6



CPGs can modulate speed, heading, and type of gait under the modulation of a few drive signals

100%


From Lecture 6

Descending modulation

Spinal cord

Reflexes

Central pattern generators



Musculoskeletal system

Respective Role
in motor control



lamprey



salamander



cat



human

"Comp ima sp

Modeling the CPG with coupled oscillators (Quadruped)

Amplitude:


$$\dot{r}_i = \alpha(\mu - r_i^2)r_i$$

Phase:


$$\dot{\theta}_i = \omega_i + \sum_{j=0}^3 r_j w_{ij} \sin(\theta_j - \theta_i - \phi_{ij})$$

Output:

$$x_{\text{foot}} = -d_{\text{step}} r_i \cos(\theta_i)$$
$$z_{\text{foot}} = \begin{cases} -h + g_c \sin(\theta_i) & \text{if } \sin(\theta_i) > 0 \\ -h + g_p \sin(\theta_i) & \text{otherwise} \end{cases}$$



Mapping CPG States to Foot Positions with Inverse Kinematics



$$\dot{r}_i = \alpha(\mu - r_i^2)r_i$$

$$\dot{\theta}_i = \omega_i$$



$$x_{\text{foot}} = -d_{\text{step}}r_i \cos(\theta_i)$$

$$z_{\text{foot}} = \begin{cases} -h + g_c \sin(\theta_i) & \text{if } \sin(\theta_i) > 0 \\ -h + g_p \sin(\theta_i) & \text{otherwise} \end{cases}$$

Gait Terminology

From Lecture 2


- *Stride duration* = the duration of a complete cycle (the period)
- *Swing phase* of a limb (period during which the limb is off the ground)
- *Stance phase* (period during which the limb touches the ground)
- *Duty factor* = Stance duration / Stride duration



Most common quadruped gaits

From Lecture 2

Classification in terms of the footfall sequences (mainly used in mathematical biology)



This project

$$\dot{r}_i = \alpha(\mu - r_i^2)r_i$$

$$\dot{\theta}_i = \omega_i + \sum_{j=0}^3 r_j w_{ij} \sin(\theta_j - \theta_i - \phi_{ij})$$

What should ϕ be for each gait?


Deep Reinforcement Learning: Review

From Lecture 7

Reinforcement Learning

An MDP is defined by:

- Set of states \mathcal{S}
- Set of actions \mathcal{A}
- Transition function $P(s' | s, a)$
- Reward function $R(s, a, s')$
- Start state s_0
- Discount factor γ
- Horizon H



- Return over a trajectory $\tau = (s_0, a_0, s_1, a_1, \dots)$

$$R(\tau) = \sum_{t=0}^{\infty} \gamma^t r_t$$

- Policy $\pi(a_t | s_t)$ maps from states s_t to actions a_t (Goal: find policy maximizing above return)
- Value function: $V^\pi(s) = \mathbb{E}_{\tau \sim \pi}[R(\tau) | s_0 = s]$
- Action-value function: $Q^\pi(s, a) = \mathbb{E}_{\tau \sim \pi}[R(\tau) | s_0 = s, a_0 = a]$
- Advantage function: $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$

From Lecture 7

Many Existing Tools for Reinforcement Learning

- RL algorithm implementations

- stable-baselines3 <https://github.com/DLR-RM/stable-baselines3>
- ray[rllib] <https://github.com/ray-project/ray>
- spinningup <https://github.com/openai/spinningup>
- tianshou <https://github.com/thu-ml/tianshou/>
- ... many others!

PPO, SAC

- Physics simulators

- pybullet <https://github.com/bulletphysics/bullet3>
- MuJoCo <https://mujoco.org>
- RaiSim <https://raisim.com>
- Isaac-Gym <https://developer.nvidia.com/isaac-gym>
- ... and others!

RL Considerations

Algorithm

- On/off policy
- Hyperparameters
- Network architecture
- Random seeds/trials

...implementation
dependent!

MDP Design Decisions

- Observation space
- Action space
- Reward function

Environment Parameters



- Simulator dynamics
- Control gains –
joint/Cartesian
- Control/environment
time step
- Noise, latency

From Lecture 7

State/Action/Reward Space: A1


s_t ? i.e.

- body (z, r, p, y)
- body velocities
- joint states



a_t ?

- motor positions/torques
- Cartesian PD



r_t ? i.e.


- body linear velocity
- energy penalty

This project: construct the MDP


From Lecture 7

Joint Position Control vs. Cartesian PD Control (PPO/SAC)



Action Space: $a_t = q_{1...N}$



Action Space: $a_t = [x_{ee_i}, y_{ee_i}, z_{ee_i}]$



How robust is your approach? To be determined at the 21.12.2021 competition



Tips

- Monitor episode length and reward mean during training
- Training should complete within 1 million timesteps for reasonable observation space, action space, and reward function choices (with no noise in the environment)
- No training on test environment (used for competition)
- Start training early!